Introduction to PlackettLuce

Heather Turner[1]**, Jacob van Etten**[2]**, David Firth**[3]**, and Ioannis Kosmidis**[4]

[1]Department of Statistics, University of Warwick, UK
[2]Bioversity International, Costa Rica
[3]Department of Statistics, University of Warwick, UK; and Alan Turing Institute, London, UK
[4]Department of Statistical Science, UCL, UK; and Alan Turing Institute, London, UK

30 November 2017

Abstract

The **PlackettLuce** package implements a generalization of the model jointly attributed to Plackett (1975) and Luce (1959) for modelling rankings data. The generalization accommodates both ties (of any order) and partial rankings (rankings of only some items). By default, the implementation adds a set of pseudo-rankings with a hypothetical item, ensuring that the network of wins and losses is always strongly connected, i.e. all items are connected to every other item by both a path of wins and a path of losses. This means that the worth of each item is always estimable with finite standard error. It also has a shrinkage effect, regularizing the estimated parameters. In addition to standard methods for model summary, **PlackettLuce** provides a method to estimate quasi-standard errors for the item parameters, so that comparison intervals can be derived even when a reference item is set. Finally the package provides a method for model-based partitioning, enabling the identification of subgroups of subjects that rank items differently.

Package

PlackettLuce 0.2.0

Contents

1   Introduction

Rankings data, in which each observation is an ordering of a set of items, arises in a range of applications, for example sports tournaments and consumer studies. A classic model for such data is the Plackett-Luce model. This model depends on Luce's axiom of choice (Luce 1959, Luce (1977)) which states that the odds of choosing item 1 over item 2 do not depend on the set of items from which the choice is made. Suppose we have a set of J items

$$S = \{i_1, i_2, \ldots, i_J\}.$$

Then under Luce's axiom, the probability of selecting some item j from S is given by

$$P(j|S) = \frac{\alpha_j}{\sum_{i \in S} \alpha_i}$$

where α_i represents the **worth** of item i. Viewing a ranking of J items as a sequence of choices — first choosing the top-ranked item from all items, then choosing the second-ranked item from the remaining items and so on — it follows that the probability of the ranking $i_1 \succ \ldots \succ i_J$ is

$$\prod_{j=1}^{J} \frac{\alpha_{i_j}}{\sum_{i \in A_j} \alpha_i}$$

where A_j is the set of alternatives $\{i_j, i_{j+1}, \ldots, i_J\}$ from which item i_j is chosen. The above model is also derived in Plackett (1975), hence the name Plackett-Luce model.

The **PlackettLuce** package implements a novel extension of the Plackett-Luce model that accommodates tied rankings, which may be applied to either full or partial rankings. Pseudo-rankings are utilised to obtain estimates in cases where the maximum likelihood estimates do not exist, or do not have finite standard errors. Methods are provided to obtain different parameterizations with corresponding standard errors or quasi-standard errors (that are independent of parameter constraints). There is also a method to work with the **psychotree** package to fit Plackett-Luce trees.

1.1   Comparison with other packages

Even though the Plackett-Luce model is a well-established method for analysing rankings, the software available to fit the model is limited. By considering each choice in the ranking as a multinomial observation, with one item observed out of a possible set, the "Poisson trick" (see, for example, Baker 1994) can be applied to express the model as a log-linear model, where the response is the count (one or zero) of each possible outcome within each choice. In theory, the model can then be fitted using standard software for generalized linear models. However there are a number of difficulties with this. Firstly, dummy variables must be set up to represent the presence or absence of each item in each choice and a factor created to identify each choice, which is a non-standard task. Secondly the factor identifying each choice will have many levels: greater than the number of rankings, for rankings of more than two objects. Thus there are many parameters to estimate and a standard function such as `glm` will be slow to fit the model, or may even fail as the corresponding model matrix will be too large to fit in memory. This issue can be circumvented by using the `gnm` function from **gnm**,

which provides an `eliminate` argument to efficiently estimate the effects of such a factor. Even then, the model-fitting may be relatively slow, given the expansion in the number of observations converting from rankings to counts. For example, the ranking {item 3 $\prec$ item 1 $\prec$ item 2} expands to two choices with five counts all together:

```
##      choice item 1 item 2 item 3 count
## [1,]      1      1      0      0     0
## [2,]      1      0      1      0     0
## [3,]      1      0      0      1     1
## [4,]      2      1      0      0     1
## [5,]      2      0      1      0     0
```

It is possible to aggregate observations of the same choice from the same set of alternatives, but the number of combinations increases quickly with the number of items.

Given the issues with applying general methods, custom algorithms and software have been developed. One approach is use Hunter's (2004) minorization-maximization (MM) algorithm to maximize the likelihood, which is equivalent to an iterative scaling algorithm; this algorithm is used by the **StatRank** package. Alternatively the likelihood of the observed data under the PlackettLuce model can be maximised directly using a generic optimisation method such the Broyden–Fletcher–Goldfarb–Shanno (BFGS) algorithm, as used by both the **pmr** and **hyper2** packages. Finally Bayesian methods can be used to either maximize the posterior distribution via an Expectation Maximization (EM) algorithm or to simulate the posterior distribution using Markov-chain Monte-Carlo (MCMC) techniques, both of which are provided by **PLMIX**. **PlackettLuce** offers both iterative scaling and generic optimization using either BFGS or a limited memory variant (L-BFGS) via the **lbfgs** package.

Even some of these specialized implementations can scale poorly with the number of items and/or the number of rankings as shown by the example timings in Table 2. Specifically `pmr::pl` becomes impractical to use with a moderate number of items (~10), while the functions from **hyper2** and **StatRank** take much longer to run with a large number (1000s) of unique rankings. **PlackettLuce** copes well with these moderately-sized data sets, though is not quite as fast as **PLMIX** when both the number of items and the number of unique rankings is large.

Table 1: **Features of example data sets from PrefLib (Mattei and Walsh 2013)**

	Rankings	Unique rankings	Items
Netflix	1256	24	4
T-shirt	30	30	11
Sushi	5000	4926	10

When the number of items is more than ten, it is more common to observe partial rankings than complete rankings. Partial rankings can be of two types: *sub-rankings*, where only a subset of items are ranked each time, and *incomplete rankings*, where the top n items are selected and the remaining items are unranked, but implicitly ranked lower than the top n. Sub-rankings are accommodated by the standard Plackett-Luce model, while incomplete rankings require an extension of the Plackett-Luce model. **PlackettLuce** handles sub-rankings only, while **PLMIX** handles incomplete rankings only and **hyper2** can handle both types. (**StatRank** supports partial rankings, but it is not clear in which form). Table 3 demonstrates using the NASCAR data from Hunter (2004) that **PlackettLuce** is more efficient than **hyper2** for modelling sub-rankings of a relatively large number of items.

Table 2: **Timings for fitting the Plackett-Luce model to data sets summarised in Table 1 using different packages**

See Appendix 5.1 for details and code.

	Time elapsed (s)				
	PlackettLuce	hyper2	PLMIX	pmr	StatRank
Netflix	0.042	0.056	0.165	0.410	0.940
T-shirt	0.024	0.168	0.074	a	7.681
Sushi	1.587	111.14	0.29	a	18.769

[a] Function fails to complete.

Table 3: **Timings for fitting the Plackett-Luce model to the NASCAR data from Hunter (2004)**

All rankings are unique.

Features of NASCAR data			Time elapsed (s)	
Rankings	Items	Items per ranking	PlackettLuce	hyper2
36	83	42-43	0.156	65.619

PlackettLuce is the only package out of those based on maximum likelihood estimation with the functionality to compute standard errors for the item parameters and thereby the facility to conduct inference on these parameters. Using **PLMIX**, inference may be based on the posterior distribution. In some cases, when the network of wins and losses is disconnected or weakly connected, the maximum likelihood estimate does not exist, or has infinite standard error; such issues are handled in **PlackettLuce** by utilising pseudo-rankings. This is similar to incorporating prior information as in the Bayesian approach.

PlackettLuce is also the only package that can accommodate tied rankings, through a novel extension of the Plackett-Luce model. On the other hand **hyper2** is currently the only package that can handle rankings of combinations of items, for example team rankings in sports. **PLMIX** offers the facility to model heterogeneous populations of subjects that have different sets of worth parameters via mixture models. This is similar in spirit to the model-based partitioning offered by **PlackettLuce**, except here the sub-populations are defined by binary splits on subject attributes. A summary of the package features is given in Table 4.

Table 4: **Features of packages for fitting the Plackett-Luce model**

Feature	PlackettLuce	hyper2	pmr	StatRank	PLMIX
Inference	Frequentist	No	No	No	Bayesian
Disconnected networks	Yes	No	No	No	Yes
Ties	Yes	No	No	No	No
Teams	No	Yes	No	No	No
Heterogenous case	Trees	No	No	No	Mixtures

2 Methods

2.1 Extended Plackett-Luce model

The **PlackettLuce** package permits rankings of the form

$$R = \{C_1, C_2, \ldots, C_J\}$$

where the items in set C_1 are ranked higher than (better than) the items in C_2, and so on. If there are multiple objects in set C_j these items are tied in the ranking. For a set S, let

$$f(S) = \delta_{|S|} \left(\prod_{i \in S} \alpha_i \right)^{\frac{1}{|S|}}$$

where $|S|$ is the cardinality of the set, δ_n is a parameter representing the prevalence of ties of order n, and α_i is a parameter representing the worth of item i. Then under an extension of the Plackett-Luce model allowing ties up to order D, the probability of the ranking R is given by

$$\prod_{j=1}^{J} \frac{f(C_j)}{\sum_{k=1}^{\min(D_j, D)} \sum_{S \in \binom{A_j}{k}} f(S)} \tag{1}$$

where D_j is the cardinality of C_j, A_j is the set of alternatives from which C_j is chosen, and $\binom{A_j}{k}$ is all the possible choices of k items from A_j. The value of D can be set to the maximum number of tied items observed in the data, so that $\delta_n = 0$ for $n > D$.

When the worth parameters are constrained to sum to one, they represent the probability that the corresponding item comes first in a ranking of all items, given that first place is not tied.

The 2-way tie prevalence parameter δ_2 is interpretable via the probability that two given items *of equal worth* tie for first place, given that the first place is not a 3-way or higher tie. Specifically, that probability is $\delta_2/(2 + \delta_2)$.

The 3-way and higher tie-prevalence parameters are interpretable similarly, in terms of tie probabilities among equal-worth items.

2.1.1 Pudding example (with ties)

When each ranking contains only two items, then the model in Equation **1** reduces to extended Bradley-Terry model proposed by Davidson (1970) for paired comparisons with ties. The `pudding` data set, available in **PlackettLuce**, provides the data from Example 2 of that paper, in which respondents were asked to test two brands of chocolate pudding from a total of six brands. For each comparison of brands i and j, the data set gives the frequencies that brand i was preferred (w_{ij}), that brand j was preferred (w_{ji}) and that the brands were tied (t_{ij}).

```
library(PlackettLuce)
head(pudding)
##   i j r_ij w_ij w_ji t_ij
```

```
## 1 1 2   57   19   22   16
## 2 1 3   47   16   19   12
## 3 2 3   48   19   19   10
## 4 1 4   54   18   23   13
## 5 2 4   51   23   19    9
## 6 3 4   54   19   20   15
```

First we create a matrix representing each unique ranking, as required by `PlackettLuce`, the model-fitting function in **PlackettLuce**

```
nr <- 3*nrow(pudding)
R <- matrix(0, nrow = nr, ncol = 6,
            dimnames = list(NULL, seq_len(6)))
i <- rep(pudding$i, 3)
j <- rep(pudding$j, 3)
R[cbind(seq_len(nr), i)] <- rep(c(1, 2, 1), each = nrow(pudding))
R[cbind(seq_len(nr), j)] <- rep(c(2, 1, 1), each = nrow(pudding))
head(R, 3)
##      1 2 3 4 5 6
## [1,] 1 2 0 0 0 0
## [2,] 1 0 2 0 0 0
## [3,] 0 1 2 0 0 0
tail(R, 3)
##       1 2 3 4 5 6
## [43,] 0 0 1 0 0 1
## [44,] 0 0 0 1 0 1
## [45,] 0 0 0 0 1 1
```

The matrix `R` represents first the wins of i over j, then the wins of j over i and finally the ties, so is three times as long as the original data. Each column represents a brand. In each row `0` represents an unranked brand (not in the comparison), `1` represents the brand(s) ranked in first place and `2` represents the brand in second place, if applicable.

This matrix can used to specify the rankings in the call to `PlackettLuce`, however there are some advantages to creating a formal `"rankings"` object, which `PlackettLuce` does internally if necessary:

```
R <- as.rankings(R)
head(R)
## [1] "1 > 2" "1 > 3" "2 > 3" "1 > 4" "2 > 4" "3 > 4"
tail(R)
## [1] "4 = 5" "1 = 6" "2 = 6" "3 = 6" "4 = 6" "5 = 6"
```

The `as.rankings` method checks that the rankings are specified as dense rankings, i.e. consecutive integers with no rank skipped for tied items, recoding as necessary; drops rankings with less than two items since these are uninformative, and adds column names if necessary. As can be seen above, the print method displays the rankings in a more readable form.

To specify the full set of rankings, we need the frequency of each ranking, which will be specified to the model-fitting function as a weight vector:

```
w <- unlist(pudding[c("w_ij", "w_ji", "t_ij")])
```

Now we can fit the model with `PlackettLuce`, passing the rankings matrix and the weight vector as arguments. Setting `npseudo = 0` means that standard maximum likelihood estimation is performed and `maxit = 7` limits the number of iterations to obtain the same worth parameters as Davidson (1970):

```
mod <- PlackettLuce(R, weights = w, npseudo = 0, maxit = 7)
## Warning in PlackettLuce(R, weights = w, npseudo = 0, maxit = 7): Iterations
## have not converged.
coef(mod, log = FALSE)
##         1         2         3         4         5         6      tie2
## 0.1389750 0.1729960 0.1618192 0.1653960 0.1587648 0.2020490 0.7466951
```

Note here we specify `log = FALSE` to give the parameterization as in Equation **1**. In the next section we discuss why it is more appropriate to use the log scale for inference.

2.2   Inference

A standard way to report model parameter estimates is to report them along with their corresponding standard error. This is an indication of the estimate's precision, however implicitly this invites comparison with zero. Such comparison is made explicit in many summary methods for models in R, with the addition of partial t or Z tests testing the null hypothesis that the parameter is equal to zero, given the other parameters in the model. However this hypothesis is generally not of interest for the worth parameters in a Plackett-Luce model: we expect most items to have *some* worth, the question is whether the items differ in their worth. In addition, a Z test based on asymptotic normality of the maximum likelihood estimate will not be appropriate for worth parameters near zero or one, since it does not take account of the fact that the parameters cannot be outside of these limits.

On the log scale however, there are no bounds on the parameters and we can set a reference level to provide meaningful comparisons. By default, the summary method for `"PlackettLuce"` objects sets the first item as the reference:

```
summary(mod)
## Call: PlackettLuce(rankings = R, npseudo = 0, weights = w, maxit = 7)
##
## Coefficients:
##       Estimate Std. Error z value Pr(>|z|)
## 1       0.0000         NA      NA       NA
## 2       0.2190     0.1872   1.170 0.242120
## 3       0.1522     0.1935   0.786 0.431598
## 4       0.1741     0.1882   0.925 0.355065
## 5       0.1331     0.1927   0.691 0.489634
## 6       0.3742     0.1924   1.945 0.051755 .
## tie2   -0.2921     0.0825  -3.541 0.000399 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual deviance:  1619.4  on  1484  degrees of freedom
## AIC:  1631.4
## Number of iterations: 8
```

None of the Z tests for the item parameters provides significant evidence against the null hypothesis, which is consistent with the test for equal preferences presented in Davidson (1970). The tie parameter is also shown on the log scale here, but it is an integral part of the model rather than a parameter of interest for inference, and its scale is not relevant.

The reference level for the item parameters can be changed via the `ref` argument, for example setting to `NULL` sets the mean worth as the reference:

```
summary(mod, ref = NULL)
## Call: PlackettLuce(rankings = R, npseudo = 0, weights = w, maxit = 7)
##
## Coefficients:
##         Estimate Std. Error z value Pr(>|z|)
## 1      -0.175426   0.121939  -1.439 0.150253
## 2       0.043548   0.121813   0.358 0.720715
## 3      -0.023240   0.126818  -0.183 0.854599
## 4      -0.001377   0.121998  -0.011 0.990994
## 5      -0.042296   0.127049  -0.333 0.739203
## 6       0.198790   0.126580   1.570 0.116305
## tie2 -0.292098   0.082500  -3.541 0.000399 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual deviance:  1619.4  on  1484  degrees of freedom
## AIC:  1631.4
## Number of iterations: 8
```

As can be seen from the output above, the standard error of the item parameters changes with the reference level. Therefore in cases where there is not a natural reference (e.g. own brand versus competitor's brands), the inference can depend on an arbitrary choice. This problem can be handled through the use of *quasi standard errors* that remain constant for a given item regardless of the reference. In addition quasi standard errors are defined for the reference item, so even in cases where there is a natural reference, the uncertainty around the worth of that item can still be represented.

Quasi standard errors for the item parameters are implemented via a method for the `qvcalc` function from the **qvcalc** package:

```
qv <- qvcalc(mod)
summary(qv)
## Model call:  PlackettLuce(rankings = R, npseudo = 0, weights = w, maxit = 7)
##        estimate         SE    quasiSE     quasiVar
##    1 0.0000000 0.0000000 0.1328834 0.01765799
##    2 0.2189744 0.1872050 0.1327324 0.01761790
##    3 0.1521861 0.1935076 0.1395693 0.01947959
##    4 0.1740489 0.1881999 0.1330195 0.01769418
##    5 0.1331303 0.1926938 0.1399200 0.01957760
##    6 0.3742162 0.1923830 0.1391874 0.01937314
## Worst relative errors in SEs of simple contrasts (%):  -0.8 0.8
## Worst relative errors over *all* contrasts (%):  -1.7 1.7
```

Again by default, the first item is taken as the reference, but this may be changed via a `ref` argument. The plot method for the returned object visualizes the item parameters (log-worth parameters) along with comparison intervals - item parameters for which the comparison intervals do not cross are significantly different:

```
plot(qv, xlab = "Brand of pudding", ylab = "Worth (log)", main = NULL)
```
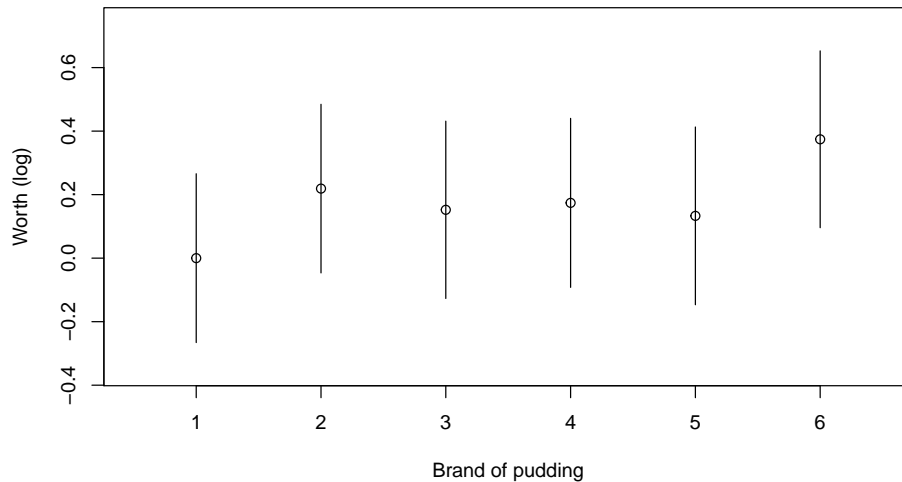

Figure 1: Worth of brands of chocolate pudding
Intervals based on quasi-standard errors.

The quasi-variances allow comparisons that are *approximately* correct, for every possible contrast among the parameters. The routine error report in the last two lines printed above by `summary(qv)` tells us that, in this example, the approximation error is always very small: the approximation error for the standard error of any simple contrast among the parameters is less than 0.8%.

2.3   Disconnected networks

The wins and losses between items can be represented as a directed network. For example, consider the following set of paired comparisons:

```
R <- matrix(c(1, 2, 0, 0,
              2, 0, 1, 0,
              1, 0, 0, 2,
              2, 1, 0, 0,
              0, 1, 2, 0), byrow = TRUE, ncol = 4,
            dimnames = list(NULL, LETTERS[1:4]))
R <- as.rankings(R)
```

The `adjacency` function from **PlacketLuce** can be used to convert these to an adjacency matrix where element (i, j) is the number of times item i is ranked higher than item j:

```
A <- adjacency(R)
A
##   A B C D
```

```
## A 0 1 0 1
## B 1 0 1 0
## C 1 0 0 0
## D 0 0 0 0
## attr(,"class")
## [1] "adjacency" "matrix"
```

Using functions from **igraph** we can visualise the corresponding network:

```
library(igraph)
net <- graph_from_adjacency_matrix(A)
plot(net, edge.arrow.size = 0.5, vertex.size = 30)
```
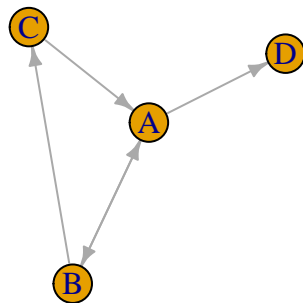

Figure 2: Network representation of toy rankings

For the worth parameters (on the log scale) to have finite maximum likelihood estimates (MLEs) and standard errors, the network must be strongly connected, i.e. there must be a path of wins and a path of losses between each pair of items. In the example above, A, B and C are strongly connected. For example, C directly loses against B and although C never directly beats B, it does beat A and A in turn beats B, so C indirectly beats B. Similar paths of wins and losses can be found for all pairs of A, B and C. On the other hand D is only observed to lose, therefore the MLE of the log-worth would be $-\infty$, with infinite standard error.

If one item always wins, the MLE of the log-worth would be $+\infty$ with infinite standard error. Or if there are clusters of items that are strongly connected with each other, but disconnected or connected only by wins or only by loses (weakly connected) to other clusters, then the maximum likelihood estimates are undefined, because there is no information on the relative worth of the clusters or one cluster is infinitely worse than the other.

The connectivity of the network can be checked with the `connectivity` function from **PlackettLuce**

```
connectivity(A)
## Network of items is not strongly connected
## $membership
## A B C D
## 1 1 1 2
##
## $csize
## [1] 3 1
##
```

```
## $no
## [1] 2
```

If the network is not strongly connected, information on the clusters within the network is returned. In this case a model could be estimated by excluding item D:

```
R2 <- R[, -4]
## Removed rankings with less than 2 items
R2
## [1] "A > B" "C > A" "B > A" "B > C"
mod <- PlackettLuce(R2, npseudo = 0)
summary(mod)
## Call: PlackettLuce(rankings = R2, npseudo = 0)
##
## Coefficients:
##   Estimate Std. Error z value Pr(>|z|)
## A   0.0000        NA      NA       NA
## B   0.8392    1.3596   0.617    0.537
## C   0.4196    1.5973   0.263    0.793
##
## Residual deviance:  5.1356  on  2  degrees of freedom
## AIC:  9.1356
## Number of iterations: 9
```

Note that since `R` is a rankings object, the rankings are automatically updated when items are dropped, so in this case the paired comparison with item D is dropped.

By default however **PlackettLuce** provides a way to handle disconnected/weakly connected networks, through the addition of pseudo-rankings. This works by adding a win and a loss between each item and a hypothetical or ghost item with fixed worth. This makes the network strongly connected so all the worth parameters are estimable. It also has an interpretation as a Bayesian prior, in particular an exchangeable prior in which all items have equal worth.

The `npseudo` argument defines the number of wins and loses with the ghost item that are added for each real item. Setting `npseudo = 0` means that no pseudo-rankings are added, so `PlackettLuce` will return the standard MLE if the network is strongly connected and throw an error otherwise. The larger `npseudo` is, the stronger the influence of the prior, by default `npseudo` is set to 0.5, so each pseudo-ranking is weighted by 0.5. This is enough to connect the network, but is a weak prior. In this toy example, the item parameters change quite considerably:

```
mod2 <- PlackettLuce(R)
coef(mod2)
##         A          B          C          D
##  0.0000000  0.5184180  0.1354701 -1.1537567
```

This is because there are only 5 rankings, so there is not much information in the data. In more realistic examples, the default prior will have a weak shrinkage effect, shrinking the items' worth parameters towards $1/N$, where N is the number of items.

For a practical example, we consider the NASCAR data from Hunter (2004). This collects the results of the 36 races in the 2002 NASCAR season in the United States. Each race involves 43 drivers out of a total of 87 drivers. The `d_nascar` data provided by **PlackettLuce** records the results as an ordering of the drivers in each race:

```
data(nascar)
nascar[1:2, 1:45]
##      rank1 rank2 rank3 rank4 rank5 rank6 rank7 rank8 rank9 rank10 rank11
## [1,]    83    18    20    48    53    51    67    72    32     42      2
## [2,]    52    72     4    82    60    31    32    66     3     44      2
##      rank12 rank13 rank14 rank15 rank16 rank17 rank18 rank19 rank20 rank21
## [1,]     31     62     13     37      6     60     66     33     77     56
## [2,]     48     83     67     41     77     33     61     45     38     51
##      rank22 rank23 rank24 rank25 rank26 rank27 rank28 rank29 rank30 rank31
## [1,]     63     55     70     14     43     71     35     12     44     79
## [2,]     14     42     62     35     12     25     37     34      6     18
##      rank32 rank33 rank34 rank35 rank36 rank37 rank38 rank39 rank40 rank41
## [1,]      3     52      4      9     45     41     61     34     39     49
## [2,]     79     39     59     43     55     49     56      9     53      7
##      rank42 rank43 rank44 rank45
## [1,]     15     82      0      0
## [2,]     13     71      0      0
```

For example, in the first race, driver 83 came first, followed by driver 18 and so on. Ranks 43 to 87 are zero for all races. We can convert these orderings to rankings using `as.rankings` with `input = "ordering"`:

```
R <- as.rankings(nascar, input = "ordering")
R[1:2,]
## [1] "83 > 18 > 20 > 48 > 53 > 51 > 67 > 7 ..."
## [2] "52 > 72 > 4 > 82 > 60 > 31 > 32 > 66 ..."
```

The names corresponding to the driver IDs are available as an attribute of `nascar`; since the columns of `R` now correspond to the drivers, we can use these as the column names:

```
colnames(R) <- attr(nascar, "drivers")
R[1:3, 1:3, as.rankings = FALSE]
##      Austin Cameron Bill Elliott Bobby Hamilton
## [1,]              0          11             32
## [2,]              0          11              9
## [3,]              0           8             43
R[1:3]
## [1] "Ward Burton > Elliott Sadler > Geoff ..."
## [2] "Matt Kenseth > Sterling Marlin > Bob ..."
## [3] "Sterling Marlin > Jeremy Mayfield >  ..."
```

Maximum likelihood estimation cannot be used in this example, because four drivers placed last in each race they entered. So Hunter (2004) dropped these four drivers to fit the Plackett-Luce model, which we can reproduce as follows:

```
keep <- seq_len(83)
R2 <- R[, keep]
mod <- PlackettLuce(R2, npseudo = 0)
```

In order to demonstrate the correspondence with the results from Hunter (2004), we order the item parameters by the driver's average rank:

```
avRank <- apply(R, 2, function(x) mean(x[x > 0]))
coefs <- round(coef(mod)[order(avRank[keep])], 2)
head(coefs, 3)
##     PJ Jones Scott Pruett  Mark Martin
##         4.15         3.62         2.08
tail(coefs, 3)
##  Dave Marcis Dick Trickle    Joe Varde
##         0.03        -0.31        -0.15
```

Now we fit the Plackett-Luce model to the full data, using the default pseudo-rankings method.

```
mod2 <- PlackettLuce(R)
```

For items that were in the previous model, we see that the log-worth parameters generally shrink towards zero:

```
coefs2 <- round(coef(mod2), 2)
coefs2[names(coefs)[1:3]]
##     PJ Jones Scott Pruett  Mark Martin
##         3.20         2.77         1.91
coefs2[names(coefs)[81:83]]
##  Dave Marcis Dick Trickle    Joe Varde
##         0.02        -0.38        -0.12
```

The new items have relative large negative log worth

```
coefs2[84:87]
## Andy Hillenburg  Gary Bradberry  Jason Hedlesky    Randy Renfrow
##          -2.17           -1.74           -1.59            -1.77
```

Nonetheless, the estimates are finite and have finite standard errors:

```
coef(summary(mod2))[84:87,]
##                   Estimate Std. Error    z value    Pr(>|z|)
## Andy Hillenburg -2.171065   1.812994 -1.1975029 0.2311106
## Gary Bradberry  -1.744754   1.855365 -0.9403829 0.3470212
## Jason Hedlesky  -1.590764   1.881708 -0.8453828 0.3978971
## Randy Renfrow   -1.768629   1.904871 -0.9284767 0.3531604
```

Note that the reference here is simply the driver that comes first alphabetically: A. Cameron. We can plot the quasi-variances for a better comparison:

```
qv <- qvcalc(mod2)
qv$qvframe <- qv$qvframe[order(coef(mod2)),]
plot(qv, xlab = NULL, ylab = "Ability (log)", main = NULL, xaxt="n")
axis(1, at = seq_len(87), labels = rownames(qv$qvframe), las = 2, cex.axis = 0.7)
```

As in the previous example, we can use `summary(qv)` to see a report on the accuracy of the quasi-variance approximation. In this example the error of approximation, across the standard errors of all of the 3741 possible simple contrasts (contrasts between pairs of the 87 driver-specific parameters), ranges between -0.7% and +6.7% — which is still remarkably
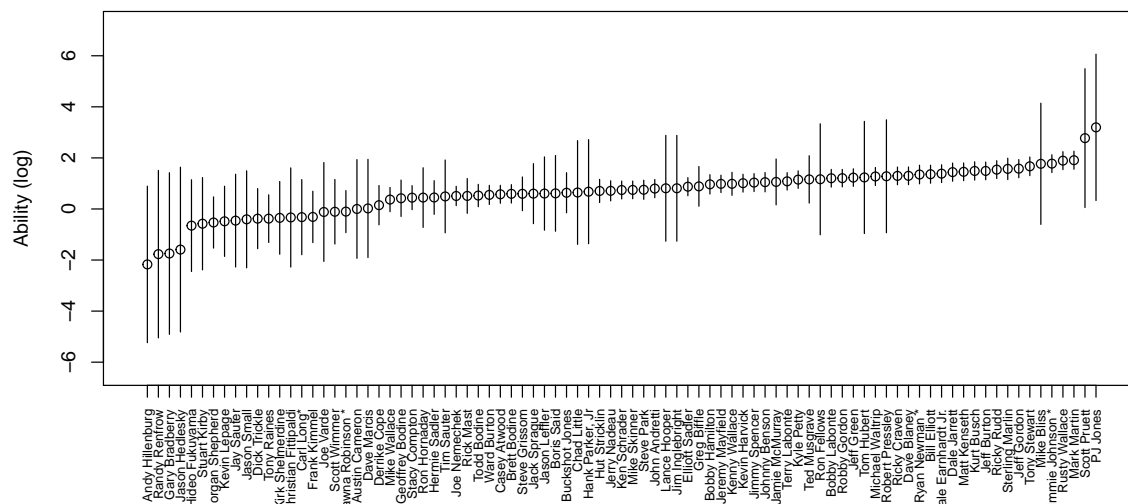
Figure 3: Ability of drivers based on NASCAR 2002 season
Intervals based on quasi-standard errors.

accurate, and which means that the plot of comparison intervals is a good visual guide to the uncertainty about drivers' relative abilities. [The results of using `summary(qv)` are not shown here, as they would occupy too much space.]

Although the prior pseudo-rankings are only necessary to add when the network is incomplete, the default behaviour is always to use them (with a weight of 0.5) because the small shrinkage effect that the pseudo-data provides also reduces both the bias and the variance of the estimates.

3   Plackett-Luce Trees

A Plackett-Luce model that assumes the same worth parameters across all rankings may sometimes be an over-simplification. For example, if rankings are made by different judges, the worth parameters may vary between judges with different characteristics. Model-based partitioning provides an automatic way to determine subgroups of the judges with significantly different sets of worth parameters, based on judge covariates. A Plackett-Luce tree is constructed via the following steps:

1. Fit a Plackett-Luce model to the full data.
2. Assess the stability of the worth parameters with respect to each available covariate.
3. If there is significant instability, split the full data by the covariate with the strongest instability and use the cut-point with the highest improvement in model fit.
4. Repeat steps 1-3 until there are no more significant instabilities, or a split produces a sub-group below a given size threshold.

This is an extension of Bradley-Terry trees, implemented in the R package **psychotree** and described in more detail by Strobl, Wickelmaier, and Zeileis (2011).

To illustrate this approach, we consider data from a trial of different varieties of bean in Nicaragua, run by Bioversity International (Van Etten et al. 2016). Farmers were asked to grow three experimental varieties of bean in one of the growing seasons, Primera (May - August), Postrera (September - October) or Apante (November - January). At the end of the

season, they were asked which variety they thought was best and which variety they thought was worse, to give a ranking of the three varieties. In addition, they were asked to compare each trial variety to the standard local variety and say whether it was better or worse.

The data are provided as the dataset `beans` in Plackett-Luce. The data require some preparation to collate the rankings, which is detailed in Appendix 5.2. The same code is provided in the examples section of the help file of `beans`

```r
example("beans", package = "PlackettLuce")
```

The result is a `rankings` object `R` with all rankings of the three experimental varieties and the output of their comparison with the local variety.

In order to fit a Plackett-Luce tree, we need to create a `grouped_rankings` object, that defines how the rankings map to the covariate values. In this case we wish to group by each record in the original data set, so we create a `"grouped_rankings"` object that identifies the record number for each of the four rankings from each farmer (one ranking of order three plus three pairwise rankings with the local variety):

```r
G <- grouped_rankings(R, rep(seq_len(nrow(beans)), 4))
format(head(G, 2), width = 50)
##                                                         1
##   "PM2 Don Rey > SJC 730-79 > BRT 103-182, Local > BRT 103-182, ..."
##                                                         2
## "INTA Centro Sur > INTA Sequia > INTA Rojo, Local > INTA Rojo, ..."
```

For each record in the original data, we have three covariates: `season` the season-year the beans were planted, `year` the year of planting, and `maxTN` the maximum temperature at night during the vegetative cycle. The following code fits a Plackett-Luce tree with up to three nodes and at least 5% of the records in each node:

```r
beans$year <- factor(beans$year)
tree <- pltree(G ~ ., data = beans[c("season", "year", "maxTN")],
               minsize = 0.05*n, maxdepth = 3)
tree
```

The algorithm identifies three nodes, with the first split defined by high night-time temperatures, and the second splitting the single Primera season from the others. So for early planting in regions where the night-time temperatures were not too high, INTA Rojo (7) was most preferred, closely followed by the local variety. During the regular growing seasons (Postrera and Apante) in regions where the night-time temperatures were not too high, the local variety was most preferred, closely followed by INTA Sequia (8). Finally in regions where the maximum night-time temperature was high, INTA Sequia (8) was most preferred, closely followed by BRT 103-182 (2) and INTA Centro Sur (3). A plot method is provided to visualise the tree:
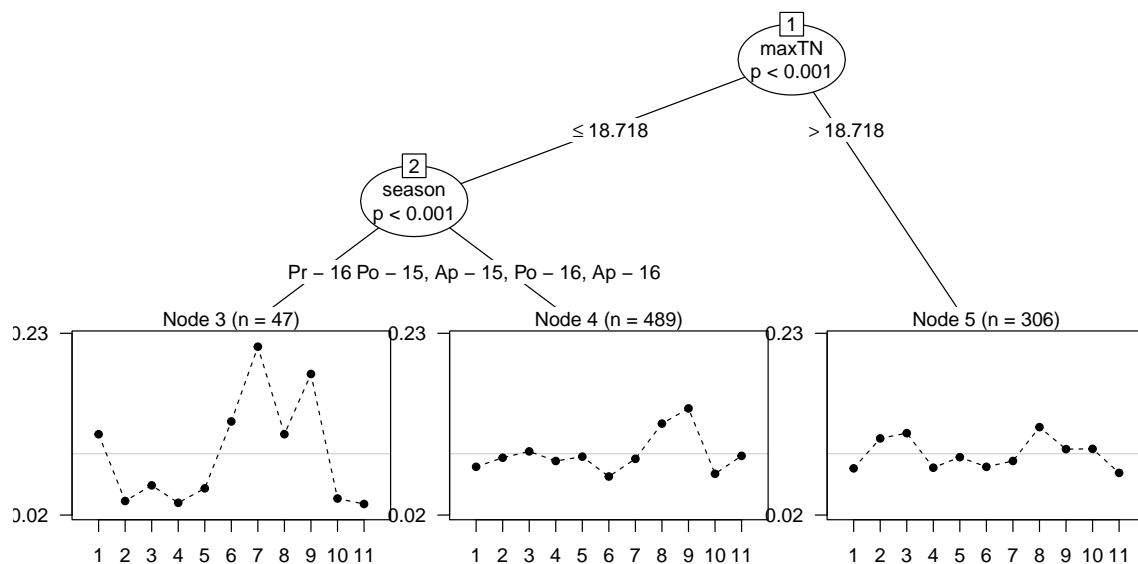
```r
plot(tree, names = FALSE, abbreviate = 2)
```

Figure 4: Worth parameters for the ten trial varieties and the local variety for each node in the Plackett-Luce tree
Varieties are 1: ALS 0532-6, 2: BRT 103-182, 3: INTA Centro Sur, 4: INTA Ferroso, 5: INTA Matagalpa, 6: INTA Precoz, 7: INTA Rojo, 8: INTA Sequia, 9: Local, 10: PM2 Don Rey, 11: SJC 730-79.

4   Discussion

PlackettLuce is a feature-rich package for the handling of ranking data. The package provides methods for importing, handling and visualising partial ranking data, and for the estimation and inference from generalizations of the Plackett-Luce model that can handle partial rankings of items and ties of arbitrary order. Disconnected item networks are handled by appropriately augmenting the data with pseudo-rankings for a hypothetical item. The package also allows for the construction of generalized Plackett-Luce trees to account for heterogeneity across the item worth parameters due to ranking-specific covariates.

Current work involves support for the online estimation from streams of partial-rankings and accounting for spatio-temporal heterogeneity in worth parameters.

5   Appendix

5.1   Timings

Data for the package comparison in Table 2 was downloaded from PrefLib (Mattei and Walsh 2013) using the `read.soc` function provided in **PlackettLuce** to read in files with the "Strict Orders - Complete List" format.

```
library(PlackettLuce)
# read in example data sets
preflib <- "http://www.preflib.org/data/election/"
netflix <- read.soc(file.path(preflib, "netflix/ED-00004-00000101.soc"))
```

```
tshirt <- read.soc(file.path(preflib, "shirt/ED-00012-00000001.soc"))
sushi <- read.soc(file.path(preflib, "sushi/ED-00014-00000001.soc"))
```

A wrapper was defined for each function in the comparison, to prepare the rankings and run
each function with reasonable defaults. The Plackett-Luce model was fitted to aggregated
rankings where possible (for `PlackettLuce`, `hyper2`, and `pmr`). Arguments were set to obtain
the maximum likelihood estimate, with the default convergence criteria. The default iterative
scaling algorithm was used for `PlackettLuce`.

```
pl <- function(dat, ...){
    # convert ordered items to ranking
    R <- as.rankings(dat[,-1], "ordering")
    # fit without adding pseudo-rankings, weight rankings by count
    PlackettLuce(R, npseudo = 0, weights = dat$n)
}
hyper2 <- function(dat, ...){
    requireNamespace("hyper2")
    # create likelihood object based on ordered items and counts
    H <- hyper2::hyper2(pnames = paste0("p", seq_len(ncol(dat) - 1)))
    for (i in seq_len(nrow(dat))){
        x <-  dat[i, -1][dat[i, -1] > 0]
        H <- H + hyper2::order_likelihood(x, times = dat[i, 1])
    }
    # find parameters to maximise likelihood
    p <- hyper2::maxp(H)
    structure(p, loglik = hyper2::lhyper2(H, p[-length(p)]))
}
plmix <- function(dat, ...){
    requireNamespace("PLMIX")
    # disaggregate data (no functionality for weights or counts)
    r <- rep(seq_len(nrow(dat)), dat$n)
    # maximum a posteriori estimate, with non-informative prior
    # K items in each ranking, single component distribution
    # default starting values do not always work so specify as uniform
    K <- ncol(dat)
    PLMIX::mapPLMIX(as.matrix(dat[r, -1]), K = K, G = 1,
                    init = list(p = rep.int(1/K, K)), plot_objective = FALSE)
}
pmr <- function(dat, ...){
    requireNamespace("pmr")
    # convert ordered items to ranking
    R <- as.rankings(dat[,-1], "ordering")
    # create data frame with counts as required by pl
    X <- as.data.frame(unclass(R))
    X$n <- dat$n
    capture.output(res <- pmr::pl(X))
    res
}
statrank <- function(dat, iter){
    requireNamespace("StatRank")
    # disaggregate data (no functionality for weights or counts)
```

```
    r <- rep(seq_len(nrow(dat)), dat$n)
    capture.output(res <- StatRank::Estimation.PL.MLE(as.matrix(dat[r, -1]),
                                                       iter = iter))
    res
}
```

When recording timings, the number of iterations for StatRank was set so that the log-likelihood on exit was equal to the log-likelihood returned by the other functions with relative tolerance 1e-6.

```
timings <- function(dat, iter = NULL,
                    fun = c("pl", "hyper2", "plmix", "pmr", "statrank")){
    res <- list()
    for (nm in c("pl", "hyper2", "plmix", "pmr", "statrank")){
        if (nm %in% fun){
            res[[nm]] <- suppressWarnings(
                system.time(do.call(nm, list(dat, iter)))[["elapsed"]])
        } else res[[nm]] <- NA
    }
    res
}
netflix_timings <- timings(netflix, 6)
tshirt_timings <- timings(tshirt, 341,
                          fun = c("pl", "hyper2", "plmix", "statrank"))
sushi_timings <- timings(sushi, 5,
                         fun = c("pl", "hyper2", "plmix", "statrank"))
```

5.2   `beans` data preparation

First we handle the best and worst rankings. These give the variety the farmer thought was best or worst, coded as A, B or C for the first, second or third variety assigned to the farmer respectively.

```
data(beans)
head(beans[c("best", "worst")], 2)
## # A tibble: 2 x 2
##    best worst
##   <chr> <chr>
## 1    C     A
## 2    B     A
```

We convert these to numeric values, allowing us to impute the middle-ranked variety (a strict ranking is assumed, so the sum of each row must be 6)

```
beans <- within(beans, {
    best <- match(best, c("A", "B", "C"))
    worst <- match(worst, c("A", "B", "C"))
    middle <- 6 - best - worst
})
head(beans[c("best", "middle", "worst")], 2)
```

```
## # A tibble: 2 x 3
##    best middle worst
##   <int>  <dbl> <int>
## 1     3      2     1
## 2     2      3     1
```

This gives an ordering of the three varieties the farmer was given. The names of these varieties are stored in separate columns

```
varieties <- as.matrix(beans[c("variety_a", "variety_b", "variety_c")])
head(varieties, 2)
##      variety_a     variety_b         variety_c
## [1,] "BRT 103-182" "SJC 730-79"      "PM2 Don Rey"
## [2,] "INTA Rojo"   "INTA Centro Sur" "INTA Sequia"
```

So we can convert the variety IDs to the variety names

```
n <- nrow(beans)
beans <- within(beans, {
    best <- varieties[cbind(seq_len(n), best)]
    worst <- varieties[cbind(seq_len(n), worst)]
    middle <- varieties[cbind(seq_len(n), middle)]
})
head(beans[c("best", "middle", "worst")], 2)
## # A tibble: 2 x 3
##             best      middle       worst
##            <chr>       <chr>       <chr>
## 1     PM2 Don Rey  SJC 730-79 BRT 103-182
## 2 INTA Centro Sur INTA Sequia   INTA Rojo
```

Next we convert these orderings to sub-rankings of the full set of varieties, including the local variety as an additional item, so that we can add the paired comparisons shortly:

```
lab <- c("Local", sort(unique(as.vector(varieties))))
R <- as.rankings(beans[c("best", "middle", "worst")],
                 input = "ordering", labels = lab)
head(R)
## [1] "PM2 Don Rey > SJC 730-79 > BRT 103-182"
## [2] "INTA Centro Sur > INTA Sequia > INTA ..."
## [3] "INTA Ferroso > INTA Matagalpa > BRT  ..."
## [4] "INTA Rojo > INTA Centro Sur > ALS 0532-6"
## [5] "PM2 Don Rey > INTA Sequia > SJC 730-79"
## [6] "ALS 0532-6 > INTA Matagalpa > INTA Rojo"
```

The comparisons with the local variety are stored in another set of columns

```
head(beans[c("var_a", "var_b", "var_c")], 2)
## # A tibble: 2 x 3
##   var_a  var_b  var_c
##   <chr>  <chr>  <chr>
## 1 Worse  Worse Better
## 2 Worse Better Better
```

The following converts each of these columns to a matrix of ordered pairs:

```
paired <- list()
for (id in c("a", "b", "c")){
    ordering <- matrix("Local", nrow = n, ncol = 2)
    worse <- beans[[paste0("var_", id)]] == "Worse"
    ## put trial variety in column 1 when it is not worse than local
    ordering[!worse, 1] <- beans[[paste0("variety_", id)]][!worse]
    ## put trial variety in column 2 when it is worse than local
    ordering[worse, 2] <- beans[[paste0("variety_", id)]][worse]
    paired[[id]] <- ordering
}
head(paired[["a"]])
##      [,1]              [,2]
## [1,] "Local"           "BRT 103-182"
## [2,] "Local"           "INTA Rojo"
## [3,] "INTA Ferroso"    "Local"
## [4,] "INTA Centro Sur" "Local"
## [5,] "INTA Sequia"     "Local"
## [6,] "ALS 0532-6"      "Local"
```

Again we convert these orderings to sub-rankings of the full set of varieties and combine them with the rankings of order three:

```
paired <- lapply(paired, as.rankings, input = "ordering", labels = lab)
R <- rbind(R, paired[["a"]], paired[["b"]], paired[["c"]])
head(R)
## [1] "PM2 Don Rey > SJC 730-79 > BRT 103-182"
## [2] "INTA Centro Sur > INTA Sequia > INTA ..."
## [3] "INTA Ferroso > INTA Matagalpa > BRT  ..."
## [4] "INTA Rojo > INTA Centro Sur > ALS 0532-6"
## [5] "PM2 Don Rey > INTA Sequia > SJC 730-79"
## [6] "ALS 0532-6 > INTA Matagalpa > INTA Rojo"
tail(R)
## [1] "INTA Sequia > Local"    "INTA Sequia > Local"
## [3] "BRT 103-182 > Local"    "Local > INTA Matagalpa"
## [5] "Local > INTA Rojo"      "Local > SJC 730-79"
```

References

Baker, Stuart G. 1994. "The multinomial-Poisson transformation." *Journal of the Royal Statistical Society. Series D (the Statistician)* 43 (4): 495–504. doi:10.2307/2348134.

Davidson, Roger R. 1970. "On extending the bradley-terry model to accommodate ties in paired comparison experiments." *Journal of the American Statistical Association* 65 (329): 317–28. doi:10.1080/01621459.1970.10481082.

Hunter, David R. 2004. "MM algorithms for generalized Bradley-Terry models." *Annals of Statistics* 32 (1): 384–406. doi:10.1214/aos/1079120141.

Luce, R. Duncan. 1959. *Individual Choice Behavior: A Theoretical Analysis*. doi:10.2307/2282347.

———. 1977. "The choice axiom after twenty years." *Journal of Mathematical Psychology* 15 (3): 215–33. doi:10.1016/0022-2496(77)90032-3.

Mattei, Nicholas, and Toby Walsh. 2013. "PrefLib: A Library of Preference Data." In *Proceedings of the 3rd International Conference on Algorithmic Decision Theory (Adt 2013)*. Lecture Notes in Artificial Intelligence. Springer. http://preflib.org.

Plackett, Robert L. 1975. "The Analysis of Permutations." *Appl. Statist* 24 (2): 193–202. doi:10.2307/2346567.

Strobl, Carolin, Florian Wickelmaier, and Achim Zeileis. 2011. "Accounting for Individual Differences in Bradley-Terry Models by Means of Recursive Partitioning." *Journal of Educational and Behavioral Statistics* 36 (2). SAGE PublicationsSage CA: Los Angeles, CA: 135——153. doi:10.3102/1076998609359791.

Van Etten, Jacob, Eskender Beza, Lluís Calderer, Kees van Duijvendijk, Carlo Fadda, Basazen Fantahun, Yosef Gebrehawaryat Kidane, et al. 2016. "First experiences with a novel farmer citizen science approach: crowdsourcing participatory variety selection through on-farm triadic comparisons of technologies (tricot)." *Experimental Agriculture*, 1–22. doi:10.1017/S0014479716000739.